

```
In [42]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, roc_auc
from sklearn.metrics import log_loss, f1_score, precision_score, recall_score
import matplotlib.pyplot as plt
from pandas_profiling import ProfileReport
import seaborn as sns
import pickle
%matplotlib inline
```

## Diabeties Dataset

```
In [124]: df = pd.read_csv("https://raw.githubusercontent.com/plotly/datasets/master/diabet")
```

```
In [126]: df
```

Out[126]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171
764	2	122	70	27	0	36.8	0.340
765	5	121	72	23	112	26.2	0.245
766	1	126	60	0	0	30.1	0.349
767	1	93	70	31	0	30.4	0.315

768 rows × 9 columns



## EDA and Feature Engineering

```
In [127]: # Profile report
df.profile_report(minimal=True)
```

Summarize dataset:	15/15 [00:00<00:00, 38.89it/s,
100%	Completed]
Generate report structure:	1/1 [00:12<00:00,
100%	12.72s/it]
Render HTML: 100%	1/1 [00:08<00:00, 8.25s/it]

**Observation:-** Features like glucose, BloodPressure, skinThickness, Insulin, BMI have zeros which, not possible for living body. These features are almost normally distributed so we can use mean to fill the zeros.

```
In [128]: df.columns # checking columns names
```

```
Out[128]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
                'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
                dtype='object')
```

```
In [129]: # replacing the zeros with mean
for col in ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']:
    df[col].replace(0, df[col].mean(), inplace=True)
```

In [130]:

ProfileReport(df,minimal=True)

Summarize dataset:	15/15 [00:00<00:00, 33.24it/s,
100%	Completed]
Generate report structure:	1/1 [3:11:09<00:00,
100%	11469.19s/it]
Render HTML: 100%	1/1 [00:35<00:00, 35.25s/it]

# Overview

## Dataset statistics

Number of variables	9
Number of observations	768
Missing cells	0
Missing cells (%)	0.0%
Total size in memory	54.1 KiB
Average record size in memory	72.2 B

## Variable types

Numeric	9
---------	---

## Alerts

Pregnancies has 111 (14.5%) zeros	Zeros
Outcome has 500 (65.1%) zeros	Zeros

## Reproduction

Analysis started	2022-09-09 13:01:31.240407
Analysis finished	2022-09-09 13:01:31.379103

Out[130]:


Observation:- Zeros are replaced. There is no null values in the dataset.

```
In [131]: df.isnull().sum()
```

```
Out[131]: Pregnancies      0
          Glucose          0
          BloodPressure    0
          SkinThickness     0
          Insulin           0
          BMI               0
          DiabetesPedigreeFunction  0
          Age              0
          Outcome          0
          dtype: int64
```

```
In [132]: ## Checking Outliers using box plots
          fig ,ax = plt.subplots(figsize = (20,20))
          sns.boxplot(data = df , ax = ax)
```

```
Out[132]: <AxesSubplot:>
```



**Observation:- There are outliers in many features of the dataset.**

```
In [133]: df.shape
```

```
Out[133]: (768, 9)
```

```
In [134]: df.columns
```

```
Out[134]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',  
                'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
               dtype='object')
```

```
In [151]: # As we have Less numbers of data rows so we are carefully choosing our quantiles  
q = df['Pregnancies'].quantile(.99)  
df_new = df[df['Pregnancies'] < q]  
  
q = df_new['Insulin'].quantile(.95)  
df_new = df_new[df_new['Insulin'] < q]  
  
q = df_new['Age'].quantile(.97)  
df_new = df_new[df_new['Age'] < q]  
  
q = df_new['BloodPressure'].quantile(.98)  
df_new = df_new[df_new['BloodPressure'] < q]  
  
q = df_new['SkinThickness'].quantile(.98)  
df_new = df_new[df_new['SkinThickness'] < q]  
  
q = df_new['BMI'].quantile(.99)  
df_new = df_new[df_new['BMI'] < q]  
  
q = df_new['Glucose'].quantile(.99)  
df_new = df_new[df_new['Glucose'] < q]
```

```
In [152]: df_new.shape
```


```
Out[152]: (648, 9)
```

```
In [153]: df_new.isnull().sum()
```

```
Out[153]: Pregnancies      0  
Glucose      0  
BloodPressure  0  
SkinThickness  0  
Insulin      0  
BMI          0  
DiabetesPedigreeFunction  0  
Age          0  
Outcome      0  
dtype: int64
```

```
In [154]: fig ,ax = plt.subplots(figsize = (20,20))  
sns.boxplot(data = df_new , ax = ax)
```

Out[154]: <AxesSubplot:>



**Observation:-** Quantile technique to remove outlier removed outliers from some columns.  
We can try standardscaler as well.

```
In [155]: df_new.isnull().sum()
```

```
Out[155]: Pregnancies      0
           Glucose          0
           BloodPressure    0
           SkinThickness     0
           Insulin           0
           BMI               0
           DiabetesPedigreeFunction  0
           Age              0
           Outcome          0
           dtype: int64
```

```
In [156]: # Segregating output column
y = df_new['Outcome']
y
```

```
Out[156]: 0      1
           1      0
           2      1
           3      0
           4      1
           ..
          762     0
          764     0
          765     0
          766     1
          767     0
           Name: Outcome, Length: 648, dtype: int64
```

```
In [157]: X = df_new.drop(columns=['Outcome'])
```



In [158]: X

Out[158]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
0	6	148.0	72.0	35.000000	79.799479	33.6	0
1	1	85.0	66.0	29.000000	79.799479	26.6	0
2	8	183.0	64.0	20.536458	79.799479	23.3	0
3	1	89.0	66.0	23.000000	94.000000	28.1	0
4	0	137.0	40.0	35.000000	168.000000	43.1	2
...	...	...	...	...	...	...	...
762	9	89.0	62.0	20.536458	79.799479	22.5	0
764	2	122.0	70.0	27.000000	79.799479	36.8	0
765	5	121.0	72.0	23.000000	112.000000	26.2	0
766	1	126.0	60.0	20.536458	79.799479	30.1	0
767	1	93.0	70.0	31.000000	79.799479	30.4	0


648 rows × 8 columns



```
In [159]: # Transformation of data using standardScaler()
scalar = StandardScaler()
X_scaled = scalar.fit_transform(X)
```

```
In [160]: # Boxplot plotting to check the dataset after standard scaler
df_new_scalar = pd.DataFrame(scalar.fit_transform(df_new))
fig ,ax  = plt.subplots(figsize = (20,20))
sns.boxplot(data = df_new_scalar , ax = ax)
```

Out[160]: <AxesSubplot:>



**Observation:- After StandardScaling data looks normalized.**

In [162]: X\_scaled

```
Out[162]: array([[ 0.78080278,  1.10205322,  0.09293203, ...,  0.30463329,
                  0.52496906,  1.8616411 ],
                [-0.84190473, -1.16279909, -0.4581753 , ..., -0.82128998,
                 -0.3490655 , -0.05334973],
                [ 1.42988579,  2.36030449, -0.64187774, ..., -1.35208237,
                 0.66747469,  0.04743926],
                ...,
                [ 0.45626128,  0.13140223,  0.09293203, ..., -0.88562845,
                 -0.68474543, -0.15413872],
                [-0.84190473,  0.31115241, -1.00928262, ..., -0.25832834,
                 -0.35539908,  1.55927413],
                [-0.84190473, -0.87519879, -0.09077041, ..., -0.21007449,
                 -0.46307   , -0.85966166]])
```

In [163]: y

```
Out[163]: 0      1
          1      0
          2      1
          3      0
          4      1
          ..
          762    0
          764    0
          765    0
          766    1
          767    0
          Name: Outcome, Length: 648, dtype: int64
```

```
In [180]: # To check Multicollinearity we use vif
def vif_score(x):
    scaler = StandardScaler()
    arr = scaler.fit_transform(x)
    return pd.DataFrame([[x.columns[i], variance_inflation_factor(arr,i)] for i in range(x.shape[1])])
```

```
In [181]: vif_score(X)
```

```
Out[181]:
```

	FEATURE	VIF_SCORE
0	Pregnancies	1.566153
1	Glucose	1.277012
2	BloodPressure	1.223138
3	SkinThickness	1.458437
4	Insulin	1.252390
5	BMI	1.520853
6	DiabetesPedigreeFunction	1.056995
7	Age	1.759953

**Observation:- No Multicollinearity among features**

### Train Test Split

```
In [187]: x_train, x_test, y_train, y_test = train_test_split(X_scaled , y , test_size = .2)
```

```
In [188]: x_train
```

```
Out[188]: array([[ -1.16644624, -1.59419952, -0.1729342 , ..., -1.70594397,
        -0.37756662, -0.65808368],
       [ -0.19282173,  0.52685263,  0.8277418 , ...,  0.43331023,
        -0.18755911,  1.25690716],
       [ -0.84190473,  0.27520238, -0.09077041, ..., -1.19123619,
        -0.76074844, -0.65808368],
       ...,
       [ -0.84190473, -0.15619806, -0.64187774, ...,  0.30463329,
         0.25895854, -1.06123964],
       [ -0.19282173, -0.04834795,  0.27663448, ..., -0.86954383,
        -1.12176271, -0.75887267],
       [ -1.16644624, -0.83924876, -0.1729342 , ...,  0.04608562,
        -0.64991072, -0.65808368]])
```

In [189]: `x_test`

```
Out[189]: array([[ -0.84190473, -0.62354854,  0.27663448, ..., -1.96329786,
        -0.98875745, -0.3557167 ],
        [  2.4035103 ,  0.92230303,  2.1136589 , ...,  0.78717183,
        -0.65624431,  1.9624301 ],
        [  2.7280518 , -1.19874912,  0.09293203, ..., -0.32266682,
        -0.52007226,  1.45848514],
        ...,
        [  0.13171978,  0.38305248, -0.09077041, ...,  0.41722561,
        -0.50107151, -0.75887267],
        [-1.16644624,  0.23925234, -0.09077041, ..., -0.69261303,
        -0.65624431,  0.45059523],
        [-0.84190473,  0.77850289, -2.29519972, ..., -0.483513 ,
        0.61047244, -0.96045065]])
```

## Logistics Regression

In [208]: *# logis object by default it uses lbfgs solver and penalty= l2*  
`logis = LogisticRegression()`

In [209]: `logis.fit(x_train,y_train )`

Out[209]: `LogisticRegression()`

In [210]: *# ALL default parameters we can see*  
`logis.get_params()`

```
Out[210]: {'C': 1.0,
          'class_weight': None,
          'dual': False,
          'fit_intercept': True,
          'intercept_scaling': 1,
          'l1_ratio': None,
          'max_iter': 100,
          'multi_class': 'auto',
          'n_jobs': None,
          'penalty': 'l2',
          'random_state': None,
          'solver': 'lbfgs',
          'tol': 0.0001,
          'verbose': 0,
          'warm_start': False}
```

**Observation:- By default, lbfgs solver and penalty=l2 also fit\_intercept=True**

In [211]: *# slopes values*  
`logis.coef_`

```
Out[211]: array([[ 0.28596086,  1.18507163, -0.09769945,  0.01972989, -0.10540602,
        0.60765528,  0.47934611,  0.22849868]])
```

```
In [212]: # Binary classes
logis.classes_
```

```
Out[212]: array([0, 1], dtype=int64)
```

```
In [219]: # Intercept of the best fit line
logis.intercept_
```

```
Out[219]: array([-1.12817895])
```

```
In [236]: # Prediction probabilty of class 0 and 1
pred_prob = logis.predict_proba([x_test[1]])
```

```
In [237]: # Let's predict just one values from test data
print("predicted value: ",logis.predict([x_test[1]]))
print(f"prediction probabilty of 0 is {pred_prob[0][0]} and of 1 is {pred_prob[0][1]}")

predicted value: [1]
prediction probabilty of 0 is 0.2726911636295054 and of 1 is 0.7273088363704946
```

```
In [239]: # Log probability
logis.predict_log_proba([x_test[1]])
```

```
Out[239]: array([[ -1.29941539,  -0.31840408]])
```

```
In [246]: # Model accuracy using x_test dataset
logis.score(x_test,y_test)
```

```
Out[246]: 0.7461538461538462
```

```
In [250]: # Making prediction on x_test
y_pred = logis.predict(x_test)
```

**Checking Accuracy,confusion\_matrix,f1\_score etc using sklearn.metrics**

```
In [266]: # function shows accuracy, confusion_matrix using sklearn.metrics
def model_metrics(y_test,y_pred):
    print("Accuracy: ",accuracy_score(y_test,y_pred))          # Model Accuracy
    print("Confusion Matrix: ",confusion_matrix(y_test,y_pred)) # confusion matrix
    print("log loss :",log_loss(y_test,y_pred))                # Log Loss
    print("precision score: ", precision_score(y_test,y_pred))  # precision score
    print("recall score : ",recall_score(y_test,y_pred))        # recall score
    print("F1 score : ",f1_score(y_test,y_pred))
model_metrics(y_test,y_pred)
```

```
Accuracy:  0.7461538461538462
Confusion Matrix:  [[80 10]
 [23 17]]
log loss : 8.767597053895035
precision score:  0.6296296296296297
recall score :  0.425
F1 score :  0.5074626865671642
```

### Checking model accuracy, confusion\_matrix using formula

```
In [265]: def model_eval(y_true,y_pred):
    tn, fp, fn, tp = confusion_matrix(y_test,y_pred).ravel()
    accuracy=(tp+tn)/(tp+tn+fp+fn)
    precision=tp/(tp+fp)
    recall=tp/(tp+fn)
    specificity=tn/(fp+tn)
    F1_Score = 2*(recall * precision) / (recall + precision)
    result={"Accuracy":accuracy,"Precision":precision,"Recall":recall,"Specificity":specificity}
    return result
model_eval(y_test,y_pred)
```


```
Out[265]: {'Accuracy': 0.7461538461538462,
 'Precision': 0.6296296296296297,
 'Recall': 0.425,
 'Specificity': 0.8888888888888888,
 'F1': 0.5074626865671642}
```

### ROC\_AUC\_SCORE AND GRAPH

```
In [328]: roc_auc_score(y_test,y_pred)
```

```
Out[328]: 0.6569444444444444
```

```
In [331]: # plot
fpr, tpr, thresholds = roc_curve(y_test,y_pred)
plt.plot(fpr, tpr, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--',label='ROC curve (area
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



In [ ]:

### Creating and testing model using Solver=liblinear, penalty='l1'

```
In [275]: logis_lib = LogisticRegression(verbose=1,solver='liblinear',penalty='l1')
```

```
In [276]: logis_lib.fit(x_train,y_train)
```

```
[LibLinear]
```

```
Out[276]: LogisticRegression(penalty='l1', solver='liblinear', verbose=1)
```



```
In [278]: # see all parameter
logis_lib.get_params()
```

```
Out[278]: {'C': 1.0,
          'class_weight': None,
          'dual': False,
          'fit_intercept': True,
          'intercept_scaling': 1,
          'l1_ratio': None,
          'max_iter': 100,
          'multi_class': 'auto',
          'n_jobs': None,
          'penalty': 'l1',
          'random_state': None,
          'solver': 'liblinear',
          'tol': 0.0001,
          'verbose': 1,
          'warm_start': False}
```

```
In [280]: # slopes
logis_lib.coef_
```

```
Out[280]: array([[ 0.27969855,  1.17511063, -0.07810108,  0.00365199, -0.08997112,
                   0.59970065,  0.46837545,  0.21484829]])
```

```
In [281]: # intercepts
logis_lib.intercept_
```

```
Out[281]: array([-1.10817926])
```

```
In [283]: y_pred_lib = logis_lib.predict(x_test)
```

```
In [284]: # Checking accuracy etc
model_metrics(y_test,y_pred_lib)
```


```
Accuracy:  0.7461538461538462
Confusion Matrix:  [[80 10]
 [23 17]]
log loss : 8.767597053895035
precision score:  0.6296296296296297
recall score :  0.425
F1 score :  0.5074626865671642
```

### **ROC\_AUC\_SCORE**

```
In [333]: roc_auc_score(y_test,y_pred_lib)
```

```
Out[333]: 0.6569444444444444
```

```
In [334]: # plot
fpr, tpr, thresholds = roc_curve(y_test,y_pred_lib)
plt.plot(fpr, tpr, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--',label='ROC curve (area = 0.66)')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



**Observation:- Model of solver='lbfgs' ,penalty='l2' AND Model of solver='liblinear',penalty='l1' have almost same performance**

**Creating and testing model with solver='saga' and penalty='elasticnet'**

```
In [316]: logis_saga = LogisticRegression(solver='saga',penalty='elasticnet',l1_ratio=0.5,fit_intercept=True)
```

```
In [317]: logis_saga.fit(x_train,y_train)
```

```
Out[317]: LogisticRegression(fit_intercept=False, l1_ratio=0.5, penalty='elasticnet', solver='saga')
```

```
In [318]: # see all parameter
logis_saga.get_params()
```

```
Out[318]: {'C': 1.0,
          'class_weight': None,
          'dual': False,
          'fit_intercept': False,
          'intercept_scaling': 1,
          'l1_ratio': 0.5,
          'max_iter': 100,
          'multi_class': 'auto',
          'n_jobs': None,
          'penalty': 'elasticnet',
          'random_state': None,
          'solver': 'saga',
          'tol': 0.0001,
          'verbose': 0,
          'warm_start': False}
```

```
In [319]: logis_saga.coef_ # slopes
```

```
Out[319]: array([[ 0.26350835,  1.11069258, -0.07539441, -0.03970395, -0.16687024,
                   0.52891693,  0.43414594,  0.16602705]])
```

```
In [320]: logis_saga.intercept_ #intercept will 0
```

```
Out[320]: array([0.])
```

```
In [321]: y_pred_saga = logis_saga.predict(x_test)
```

```
In [322]: model_metrics(y_test,y_pred_saga)
```

```
Accuracy: 0.6923076923076923
Confusion Matrix: [[60 30]
                   [10 30]]
log loss : 10.627500336302564
precision score: 0.5
recall score : 0.75
F1 score : 0.6
```

**Observation:- model solver='saga',penalty='elasticnet' and fit\_intercept=False has low accuracy than above models**

### AUC\_ROC\_SCORE

```
In [335]: roc_auc_score(y_test,y_pred_saga)
```

```
Out[335]: 0.7083333333333334
```

```
In [337]: fpr, tpr, thresholds = roc_curve(y_test,y_pred_saga)
plt.plot(fpr, tpr, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--',label='ROC curve (area = 0.66)')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```

